

Modern Cryptography

600.442

Lecture #16

Dr. Christopher Pappacena

Fall 2013

Public-Key Encryption

A *public-key encryption scheme* $\Pi = (\text{Gen}, \text{Enc}, \text{Dec})$ is given by:

- $\text{Gen}(n)$ outputs a *public key* pk and a *secret key* sk .
- Enc takes as input the public key and a message m and returns a ciphertext c . We write $c \leftarrow \text{Enc}_{\text{pk}}(m)$.
- Dec takes as input the secret key and a ciphertext c and returns either a message m or a decryption failure \perp . We write $m = \text{Dec}_{\text{sk}}(c)$.

We require that $\Pr[\text{Dec}_{\text{sk}}(\text{Enc}_{\text{pk}}(m)) \neq m] = \text{negl}(n)$ for some negligible function of n .

Remarks

- The public key is made available so that *anyone* can encrypt messages.
- We allow a negligible probability of decryption failure.
- We generally need a method for encoding messages as elements of a group. For RSA, we can encode a nonzero message of $\leq n - 1$ bits by viewing it as an integer in $\{1, \dots, N - 1\}$.
- For El Gamal encryption this encoding is more complicated.

The Indistinguishability Experiment

- Gen is run to produce pk and sk .
- \mathcal{A} is given pk and produces two messages m_0, m_1 of the same length.
- $b \leftarrow \{0, 1\}$ is chosen at random and \mathcal{A} is given $c \leftarrow \text{Enc}_{\text{pk}}(m_b)$.
- \mathcal{A} outputs a bit b' . We write $\text{PubK}_{\mathcal{A}, \Pi}^{\text{eav}}(n) = 1$ if $b = b'$ and say that \mathcal{A} *succeeds*; otherwise \mathcal{A} *fails*.

Definition: A public-key encryption scheme Π has indistinguishable encryptions in the presence of an eavesdropper if, for all PPT adversaries \mathcal{A} , we have

$$\Pr[\text{PubK}_{\mathcal{A}, \Pi}^{\text{eav}}(n) = 1] \leq \frac{1}{2} + \text{negl}(n)$$

for some negligible function negl .

Remark: There is *no* notion of “perfect secrecy” for a public-key encryption schemes. A computationally unbounded adversary can recover the message m from c with probability 1 (HW Problem).

Public-Key Encryption and CPA Security

In the public-key setting, the adversary is given access to the public key pk . This means two things:

- Enc *must* be randomized. Otherwise, \mathcal{A} just computes $c_b = \text{Enc}_{\text{pk}}(m_b)$ for $b = 0, 1$ and easily succeeds.
- \mathcal{A} has access to an encryption oracle and so is able to mount a chosen-plaintext attack.

Proposition: If a public-key encryption scheme Π has indistinguishable encryptions in the presence of an eavesdropper, then Π is also CPA secure.

Multiple Encryptions

We can define an experiment $\text{PubK}_{\mathcal{A}, \Pi}^{\text{mult}}(n)$, where the adversary \mathcal{A} produces a pair of message vectors $M_0 = (m_0^1, \dots, m_0^t)$ and $M_1 = (m_1^1, \dots, m_1^t)$ for some t .

The game is played the same way, with \mathcal{A} receiving a ciphertext vector C_b . \mathcal{A} returns b' and succeeds if $b = b'$.

We define security in the presence of multiple encryptions in the obvious way.

One For All and All For One

Theorem: If a public key scheme Π has indistinguishable encryptions in the presence of an eavesdropper, then it has indistinguishable multiple encryptions in the presence of an eavesdropper.

The proof uses a hybrid argument.

Proof of Theorem

In the experiment $\text{PubK}_{\mathcal{A}, \Pi}^{\text{mult}}(n)$, \mathcal{A} selects two plaintext vectors M_0 and M_1 of length t .

For $0 \leq i \leq t$, define the ciphertext vector $C^{(i)}$ by

$$C^{(i)} = (\text{Enc}_{\text{pk}}(m_0^1), \dots, \text{Enc}_{\text{pk}}(m_0^i), \text{Enc}_{\text{pk}}(m_1^{i+1}), \dots, \text{Enc}_{\text{pk}}(m_1^t)).$$

In words, $C^{(i)}$ is an encryption of the first i plaintexts of M_0 followed by the last $t - i$ plaintexts of M_1 .

As we range over all possible randomizations of Enc , each $C^{(i)}$ defines a *distribution* on vectors with t ciphertexts.

Proof of Theorem, II

Now let \mathcal{A}' be an adversary which uses \mathcal{A} as a subroutine:

- \mathcal{A}' gives pk to \mathcal{A} and receives M_0 and M_1 .
- \mathcal{A}' chooses $i \leftarrow \{1, \dots, t\}$, outputs m_0^i, m_1^i , and receives c^i .
- \mathcal{A}' computes $c^j \leftarrow \text{Enc}_{\text{pk}}(m_0^j)$ for $j < i$ and $c^j \leftarrow \text{Enc}_{\text{pk}}(m_1^j)$ for $j > i$.
- \mathcal{A}' gives (c^1, \dots, c^t) to \mathcal{A} and returns the bit b' returned by \mathcal{A} .

Proof of Theorem, III

If $b = 0$, then the ciphertext that \mathcal{A}' gives \mathcal{A} is $C^{(i)}$. We have:

$$\begin{aligned}\Pr[\mathcal{A}'(n) = 0 | b = 0] &= \sum_{j=1}^t \Pr[\mathcal{A}'(n) = 0 | b = 0 \wedge i = j] \times \Pr[i = j] \\ &= \frac{1}{t} \sum_{j=1}^t \Pr[\mathcal{A}(C^{(j)}) = 0].\end{aligned}$$

If $b = 1$, then the ciphertext is $C^{(i-1)}$ and

$$\Pr[\mathcal{A}'(n) = 1 | b = 1] = \frac{1}{t} \sum_{j=0}^{t-1} \Pr[\mathcal{A}(C^{(j)}) = 1].$$

Proof of Theorem, IV

Combining these gives

$$\begin{aligned}\Pr[\text{PubK}_{\mathcal{A}', \Pi}^{\text{eav}}(n) = 1] &= \frac{1}{2} \Pr[\mathcal{A}'(n) = 0 | b = 0] + \frac{1}{2} \Pr[\mathcal{A}'(n) = 1 | b = 1] \\ &= \frac{1}{2t} \left(\sum_{j=1}^t \Pr[\mathcal{A}(C^{(j)}) = 0] + \sum_{j=0}^{t-1} \Pr[\mathcal{A}(C^{(j)}) = 1] \right) \\ &= \frac{t-1}{2t} + \frac{1}{2t} \Pr[\text{PubK}_{\mathcal{A}, \Pi}^{\text{mult}}(n) = 1].\end{aligned}$$

If the advantages of \mathcal{A} and \mathcal{A}' are $\epsilon(n)$ and $\epsilon'(n)$ then this gives

$$\epsilon(n) = t \cdot \epsilon'(n).$$

Recap

- Indistinguishable encryptions in the presence of an eavesdropper implies indistinguishable *multiple* encryptions in the presence of an eavesdropper.
- We can bootstrap a fixed-length public-key system into one for arbitrary-length messages.
- Indistinguishability implies CPA-security for public-key encryption. As a result, any secure public-key encryption scheme must have *randomized* encryptions.
- “Textbook” RSA doesn’t use randomization, so is insecure!

Padded RSA

Let $\ell(n)$ be a function with $\ell(n) \leq 2n - 2$. To encrypt $m \in \{0, 1\}^{\ell(n)}$, choose a random $r \leftarrow \{0, 1\}^{\|N\| - \ell(n) - 1}$ and set

$$c = (r||m)^e \pmod{N}.$$

To decrypt, let $\tilde{m} = c^d \pmod{N}$ and set m equal to the low $\ell(n)$ bits of \tilde{m} .

Theorem: If $\ell(n) = O(\log n)$ and the RSA problem is hard relative to GenRSA, then this gives a CPA-secure public-key encryption scheme.

The proof uses the fact that the low-order bits give hard-core predicates for RSA.

PKCS #1 v1.5

Let k be the size of N in bytes. Messages m range from 1 to $k - 11$ bytes long.

The encryption of an s -byte message m is

$$(00000000||00000010||r||00000000||m)^e \pmod{N},$$

where r is a random string of $k - s - 3$ nonzero bytes.

This is believed to be CPA-secure but no proof is known based on the RSA assumption. It is known to *not* be CCA-secure.

El Gamal Encryption

In 1985, El Gamal introduced a public-key encryption scheme whose security is based on the DDH problem.

Let \mathcal{G} be a group generation algorithm for which the DDH problem is hard.

- Run $\mathcal{G}(n)$ to obtain (G, q, g) . Choose $x \leftarrow \mathbb{Z}_q$ and set $h = g^x$. Set $\text{pk} = (G, q, g, h)$ and $\text{sk} = (G, q, g, x)$.
- Given a message $m \in G$, choose $y \leftarrow \mathbb{Z}_q$ and set $c = (g^y, h^y m)$.
- Given $c = (c_1, c_2)$, decrypt by setting $m = c_2 / c_1^x$ ($= c_2 \circ c_1^{-x}$).

Theorem: If the DDH problem is hard relative to \mathcal{G} , then El Gamal encryption scheme Π has indistinguishable encryptions in the presence of an eavesdropper and is CPA-secure.

Proof: Let \mathcal{A} be an adversary who can break Π with advantage $\epsilon(n)$.

Define $\widetilde{\text{Enc}}_{\text{pk}}(m)$ to be $(g^y, g^z m)$ for random $y, z \leftarrow \mathbb{Z}_q$.

If $\widetilde{\text{Enc}}_{\text{pk}}$ is used in place of Enc_{pk} then \mathcal{A} will succeed with probability $1/2$ since no information about m is revealed.

Proof, Continued

We design a distinguisher D for the DDH problem. Recall that D is given (G, q, g, g^x, g^y, h) and needs to decide if $h = g^{xy}$.

- D gives $\text{pk} = (G, q, g, g^x)$ to \mathcal{A} and receives m_0, m_1 .
- D chooses b , gives $c = (g^y, hm_b)$ to \mathcal{A} , and receives b' .
- If $b' = b$ then D returns 1, otherwise D returns 0.

Proof, Continued

If $h \neq g^{xy}$, then \mathcal{A} returns 0 and 1 with probability $1/2$ each. So D returns 1 with probability $1/2$ in this case.

If $h = g^{xy}$, then \mathcal{A} returns b with probability $1/2 + \epsilon(n)$ and so D returns 1 with probability $1/2 + \epsilon(n)$ in this case.

Hence,

$$|\Pr[D(h = g^{xy}) = 1] - \Pr[D(h \neq g^{xy}) = 1]| = \epsilon(n).$$

Since the DDH problem is hard for \mathcal{G} , $\epsilon(n)$ is negligible.

Factoring and One-Way Functions

Let $t(n)$ be the maximum number of random bits needed by $\text{GenMod}(n)$ to produce (p, q, N) .

We define a function $f : \{0, 1\}^{t(n)} \rightarrow \{0, 1\}^{2n}$ as follows:

- Run $\text{GenMod}(n)$ using input $x \in \{0, 1\}^{t(n)}$ in place of the random bits.
- Given the output (p, q, N) of GenMod , let $f(x) = N$.

If f can be inverted, then the input to f can be used to find the factors of N . So if factoring is hard relative to GenMod , then f is a one-way function.

The RSA Problem and One-Way Permutations

Let $\text{GenRSA}(n) = (N, e, d)$. Then the map

$$f_{e,N} : x \mapsto x^e \pmod{N}$$

is a *permutation* of the set \mathbb{Z}_N^* , with inverse $f_{d,N}$.

The problem of computing x given $y = f_{e,N}(x)$ is exactly the RSA problem with input y .

So if the RSA problem is hard relative to GenRSA, then $f_{e,N}$ is (morally) a one-way permutation.

In order to make this idea fit the general framework of one-way functions we have to be a bit fussy. Details are in the book.

Hash Functions

Define a fixed-length hash function H as follows:

- Run $\mathcal{G}(n)$ to obtain (G, q, g) , select $h \leftarrow G$, and set $s = (G, q, g, h)$.
- Given input $x = (x_1, x_2) \in \mathbb{Z}_q \times \mathbb{Z}_q$, output $H^s(x) = g^{x_1}h^{x_2}$.

H looks like a hash function since it reduces two elements of a cyclic group of order q to a single element of a cyclic group of order q .

For it to actually compress its output we need elements of G to be encoded by no more than $2n - 2$ bits, or use a *randomness extractor*.

The Discrete Logarithm Problem and Hash Functions

Theorem: If the discrete logarithm problem is hard relative to \mathcal{G} and H compresses its input, then H is a collision-resistant hash function.

Proof: Suppose \mathcal{A} can invert H . Given $s = (G, q, g, h)$, call \mathcal{A} as a subroutine with seed s and receive x, x' .

Write $x = (x_1, x_2)$ and $x' = (x'_1, x'_2)$ and set $y = (x_1 - x'_1)/(x_2 - x'_2) \pmod{q}$.

If $H^s(x) = H^s(x')$, then $y = \log_g h$.

The above can be turned into a formal security proof – see Theorem 7.73 of Katz and Lindell.